

# project-source-code

September 7, 2020

## 1 Project Overview

- This project will focus on addressing the problem of sentiment analysis using BERT Deep Learning technique.
- BERT is a large-scale transformer-based Language Model that can be fine-tuned for different tasks.
- Original paper of BERT can be found [here](#).
- For the experiments, we will use [SMILE Twitter Emotion dataset](#).

```
[1]: import torch
import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot
from tqdm.notebook import tqdm
%matplotlib inline
```

## 2 Data Exploration

```
[2]: df = pd.read_csv("smile-annotations-final.csv", names=['id', 'text', 'category'])
df.head()
```

```
[2]:
```

	id	text	\
0	611857364396965889	@aandraous @britishmuseum @AndrewsAntonio Merc...	
1	614484565059596288	Dorian Gray with Rainbow Scarf #LoveWins (from...	
2	614746522043973632	@SelectShowcase @Tate_StIves ... Replace with ...	
3	614877582664835073	@Sofabsports thank you for following me back. ...	
4	611932373039644672	@britishmuseum @TudorHistory What a beautiful ...	

  

	category
0	nocode
1	happy
2	happy
3	happy
4	happy

As we know that id is unique for every tweet, we would like to set index based on id number of every sample.

```
[3]: df.set_index('id', inplace=True)
df.head()
```

```
[3]:
```

	text	category
id		
611857364396965889	@aandraous @britishmuseum @AndrewsAntonio Merc...	nocode
614484565059596288	Dorian Gray with Rainbow Scarf #LoveWins (from...	happy
614746522043973632	@SelectShowcase @Tate_StIves ... Replace with ...	happy
614877582664835073	@Sofabsports thank you for following me back. ...	happy
611932373039644672	@britishmuseum @TudorHistory What a beautiful ...	happy

```
[4]: df.text.iloc[100]
```

```
[4]: "* @hist_astro @britishmuseum It looks like there's some #ArtisticLicence
involved in that sketch of #SummerSolstice #sunrise at #Stonehenge."
```

```
[5]: df.category.value_counts()
```

```
[5]: nocode          1572
happy            1137
not-relevant     214
angry             57
surprise         35
sad              32
happy|surprise   11
happy|sad        9
disgust|angry    7
disgust          6
sad|disgust      2
sad|angry        2
sad|disgust|angry 1
Name: category, dtype: int64
```

```
[6]: df = df[~df.category.str.contains('\|')]
df.category.value_counts()
```

```
[6]: nocode          1572
happy            1137
not-relevant     214
angry             57
surprise         35
sad              32
disgust          6
Name: category, dtype: int64
```

```
[7]: df = df[df.category!='nocode']
df.category.value_counts()
```

```
[7]: happy          1137
     not-relevant   214
     angry          57
     surprise       35
     sad            32
     disgust        6
     Name: category, dtype: int64
```

```
[8]: labels = df.category.unique()
     labels_di = {}

     for index, label in enumerate(labels):
         labels_di[label] = index
     print(labels_di)
```

```
{'happy': 0, 'not-relevant': 1, 'angry': 2, 'disgust': 3, 'sad': 4, 'surprise': 5}
```

```
[9]: df['label'] = df.category.replace(labels_di)
     print(df.label.value_counts())
     print(df.category.value_counts())
```

```
0      1137
1       214
2        57
5        35
4        32
3         6
     Name: label, dtype: int64
happy          1137
not-relevant   214
angry          57
surprise       35
sad            32
disgust        6
     Name: category, dtype: int64
```

### 3 Data Separation

Since the dataset for the experiment is not balanced, we do not want to use simple train-test-split for this data. Because, a class with limited datasamples might not be represented in training or validation sets, which will cause problem of generalizability of the model. Therefore, we would like to use stratified split that ensures a certain portion of the examples will be in training and test splits for each class.

```
[10]: from sklearn.model_selection import train_test_split
     X_train, X_val, y_train, y_val = train_test_split(df.index.values, df.label.
     →values,
```

```

stratify=df.label.values,
→test_size=0.15,
shuffle=True,
→random_state=2020)

```

Now, we can check wheter the samples in the data were properly distributed into train and validation sets. For this, we can create another column in the dataframe, called 'data\_type', which shows wheter the sample is in training or validation set.

```

[11]: df['data_type'] = ['not_set'] * df.shape[0]
df.head()

```

```

[11]:
text \
id
614484565059596288 Dorian Gray with Rainbow Scarf #LoveWins (from...
614746522043973632 @SelectShowcase @Tate_StIves ... Replace with ...
614877582664835073 @Sofabsports thank you for following me back. ...
611932373039644672 @britishmuseum @TudorHistory What a beautiful ...
611570404268883969 @NationalGallery @ThePoldarkian I have always ...

category label data_type
id
614484565059596288 happy 0 not_set
614746522043973632 happy 0 not_set
614877582664835073 happy 0 not_set
611932373039644672 happy 0 not_set
611570404268883969 happy 0 not_set

```

```

[12]: df.loc[X_train, 'data_type'] = 'train'
df.loc[X_val, 'data_type'] = 'val'
df.head()

```

```

[12]:
text \
id
614484565059596288 Dorian Gray with Rainbow Scarf #LoveWins (from...
614746522043973632 @SelectShowcase @Tate_StIves ... Replace with ...
614877582664835073 @Sofabsports thank you for following me back. ...
611932373039644672 @britishmuseum @TudorHistory What a beautiful ...
611570404268883969 @NationalGallery @ThePoldarkian I have always ...

category label data_type
id
614484565059596288 happy 0 val
614746522043973632 happy 0 train
614877582664835073 happy 0 train
611932373039644672 happy 0 train
611570404268883969 happy 0 train

```

```

[13]: df.groupby(['category', 'label', 'data_type']).count()

```

```
[13]:
```

			text
category	label	data_type	
angry	2	train	48
		val	9
disgust	3	train	5
		val	1
happy	0	train	966
		val	171
not-relevant	1	train	182
		val	32
sad	4	train	27
		val	5
surprise	5	train	30
		val	5

## 4 Tokenizing and Encoding the data

```
[14]: from transformers import BertTokenizer
from torch.utils.data import TensorDataset
```

D:\Anaconda3\lib\site-packages\h5py\\_\_init\_\_.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.

```
from ._conv import register_converters as _register_converters
```

```
[15]: tokenizer = BertTokenizer.from_pretrained(
        'bert-base-uncased',
        do_lower_case=True)
```

```
[16]: encoded_train = tokenizer.batch_encode_plus(
        df[df['data_type']=='train'].text.values,
        add_special_tokens=True,
        return_attention_mask=True,
        pad_to_max_length=True,
        max_length=256, return_tensors='pt')

encoded_val = tokenizer.batch_encode_plus(
        df[df['data_type']=='val'].text.values,
        add_special_tokens=True,
        return_attention_mask=True,
        pad_to_max_length=True,
        max_length=256, return_tensors='pt')
print(f"{encoded_train.keys()}\n{encoded_val.keys()}")
```

Truncation was not explicitly activated but `max\_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max

length. Defaulting to 'longest\_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

D:\Anaconda3\lib\site-packages\transformers\tokenization\_utils\_base.py:1770:

FutureWarning: The `pad\_to\_max\_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max\_length'` to pad to a max length. In this case, you can give a specific length with `max\_length` (e.g. `max\_length=45`) or leave max\_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).

FutureWarning,

Truncation was not explicitly activated but `max\_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest\_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

```
dict_keys(['input_ids', 'token_type_ids', 'attention_mask'])
```

```
dict_keys(['input_ids', 'token_type_ids', 'attention_mask'])
```

```
[17]: input_ids_train = encoded_train['input_ids']
      attention_masks_train = encoded_train['attention_mask']
      labels_train = torch.tensor(df[df['data_type'] == 'train'].label.values)

      input_ids_val = encoded_val['input_ids']
      attention_masks_val = encoded_val['attention_mask']
      labels_val = torch.tensor(df[df['data_type'] == 'val'].label.values)
```

```
[18]: dataset_train = TensorDataset(input_ids_train,
                                   attention_masks_train,
                                   labels_train)

      dataset_val = TensorDataset(input_ids_val,
                                   attention_masks_val,
                                   labels_val)

      print(f"There are {len(dataset_train)} samples in the training data!")
      print(f"There are {len(dataset_val)} samples in the validation data!")
```

There are 1258 samples in the training data!

There are 223 samples in the validation data!

## 5 Formulating a Model

```
[19]: from transformers import BertForSequenceClassification
      model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
                                                            num_labels = len(labels_di),
                                                            output_attentions=False,
```

```
output_hidden_states=False)
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForSequenceClassification: ['cls.predictions.bias',

```
'cls.predictions.transform.dense.weight',  
'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight',  
'cls.seq_relationship.weight', 'cls.seq_relationship.bias',  
'cls.predictions.transform.LayerNorm.weight',  
'cls.predictions.transform.LayerNorm.bias']
```

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPretraining model).

- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:

```
['classifier.weight', 'classifier.bias']
```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

## 6 Creating DataLoaders

```
[20]: from torch.utils.data import DataLoader, RandomSampler, SequentialSampler
```

```
bs = 8  
dataloader_train = DataLoader(dataset_train,  
                              sampler=RandomSampler(dataset_train),  
                              batch_size=bs)  
dataloader_val = DataLoader(dataset_val,  
                             sampler=RandomSampler(dataset_val),  
                             batch_size=bs*4)
```

## 7 Setting an optimizer and scheduler

```
[21]: from transformers import AdamW, get_linear_schedule_with_warmup
```

```
epochs = 10  
optimizer = AdamW(model.parameters(),  
                  lr=1e-5, eps=1e-8)  
scheduler = get_linear_schedule_with_warmup(optimizer,  
                                             num_warmup_steps=0,
```

```
↳ num_training_steps=len(dataloader_train)*epochs
```

## 8 Defining evaluation metrics

Since we have imbalanced dataset, we may be interested in using f1-score, which is one of the most appropriate evaluation metrics out there to be used for imbalanced data.

```
[22]: from sklearn.metrics import f1_score

def f1_score_func(preds, targs):
    preds_flat = np.argmax(preds, axis=1).flatten()
    targs_flat = targs.flatten()
    return f1_score(targs_flat, preds_flat, average='weighted')

def accuracy_per_class(preds, labels):

    label_di_inverse = {v:k for k, v in labels_di.items()}
    preds_flat = np.argmax(preds, axis=1).flatten()
    labels_flat = labels.flatten()

    for label in np.unique(labels_flat):
        y_preds = preds_flat[labels_flat==label]
        y_true = labels_flat[labels_flat==label]
        print(f"Class: {label_di_inverse[label]}")
        print(f"Accuracy: {len(y_preds[y_preds==label])}/{len(y_true)}\n")
```

## 9 Training a BERT model

```
[23]: import random

seed_val = 17
random.seed(seed_val)
np.random.seed(seed_val)
torch.manual_seed(seed_val)
torch.cuda.manual_seed_all(seed_val)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)
model.to(device)
```

cuda

```
[23]: BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
```



```

(word_embeddings): Embedding(30522, 768, padding_idx=0)
(position_embeddings): Embedding(512, 768)
(token_type_embeddings): Embedding(2, 768)
(LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
(dropout): Dropout(p=0.1, inplace=False)
)
(encoder): BertEncoder(
  (layer): ModuleList(
    (0): BertLayer(
      (attention): BertAttention(
        (self): BertSelfAttention(
          (query): Linear(in_features=768, out_features=768, bias=True)
          (key): Linear(in_features=768, out_features=768, bias=True)
          (value): Linear(in_features=768, out_features=768, bias=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
          (dense): Linear(in_features=768, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
    )
    (intermediate): BertIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
    (output): BertOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
)
  (1): BertLayer(
    (attention): BertAttention(
      (self): BertSelfAttention(
        (query): Linear(in_features=768, out_features=768, bias=True)
        (key): Linear(in_features=768, out_features=768, bias=True)
        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
      (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
    (intermediate): BertIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
    )
  )
)

```

```

    )
    (output): BertOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
(2): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(3): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
)

```

```

(output): BertOutput(
  (dense): Linear(in_features=3072, out_features=768, bias=True)
  (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
  (dropout): Dropout(p=0.1, inplace=False)
)
)
(4): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(5): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(

```

```

        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(6): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(7): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)

```

```

        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(8): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(9): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)

```

```

        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(10): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(11): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)

```

```

    )
    )
    )
    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (dropout): Dropout(p=0.1, inplace=False)
  (classifier): Linear(in_features=768, out_features=6, bias=True)
)

```

```

[24]: def evaluate(dataloader_val):

    model.eval()

    loss_val_total = 0
    preds, true_vals = [], []

    for batch in tqdm(dataloader_val):

        batch = tuple(b.to(device) for b in batch)

        inputs = {'input_ids'      : batch[0],
                  'attention_mask': batch[1],
                  'labels'         : batch[2]}

        with torch.no_grad():

            outputs = model(**inputs)
            loss = outputs[0]
            logits = outputs[1]
            loss_val_total += loss.item()

            logits = logits.detach().cpu().numpy()
            label_ids = inputs['labels'].cpu().numpy()
            preds.append(logits)
            true_vals.append(label_ids)

    loss_val_avg = loss_val_total / len(dataloader_val)

    preds = np.concatenate(preds, axis=0)
    true_vals = np.concatenate(true_vals, axis=0)

    return loss_val_avg, preds, true_vals

```

```
[25]: for epoch in tqdm(range(1, epochs+1)):

    model.train()
    loss_train_total = 0
    progress_bar = tqdm(dataloader_train,
                        desc='Epoch {:1d}'.format(epoch),
                        leave=False, disable=False)

    for batch in progress_bar:
        model.zero_grad()
        batch = tuple(b.to(device) for b in batch)

        inputs = {'input_ids'      : batch[0],
                  'attention_mask': batch[1],
                  'labels'         : batch[2]}

        outputs = model(**inputs)

        loss = outputs[0]
        loss_train_total += loss.item()
        loss.backward()

        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
        optimizer.step()
        scheduler.step()

        progress_bar.set_postfix({'Training loss': '{:.3f}'.format(loss.item()/
→len(batch))})

#     torch.save(model.state_dict(), f'Models/BERT_ft_epoch{epoch}.model')
    tqdm.write(f'\nEpoch {epoch}')
    loss_train_avg = loss_train_total/len(dataloader_train)
    tqdm.write(f"Training loss: {loss_train_avg}")

    val_loss, preds, true_vals = evaluate(dataloader_val)
    val_f1 = f1_score_func(preds, true_vals)
    tqdm.write(f"Validation loss: {val_loss}")
    tqdm.write(f"F1 Score: {val_f1}")
```

```
HBox(children=(FloatProgress(value=0.0, max=10.0), HTML(value='')))
```

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 1', max=158.0, style=ProgressStyle(
```

Epoch 1

Training loss: 0.8713658031406282



```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value='')))
```

Validation loss: 0.6209014100687844

F1 Score: 0.6953185953656175

D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:

UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn\_for)

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 2', max=158.0, style=ProgressStyle(
```

Epoch 2

Training loss: 0.5248276907243307

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value='')))
```

Validation loss: 0.4816068027700697

F1 Score: 0.8275480263522117

D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:

UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn\_for)

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 3', max=158.0, style=ProgressStyle(
```

Epoch 3

Training loss: 0.38210659294943267

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value='')))
```

Validation loss: 0.39346700268132345

F1 Score: 0.8527610995571987

D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:

UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn\_for)

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 4', max=158.0, style=ProgressStyle(
```

Epoch 4

Training loss: 0.25785991115660606

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value=''))))
```

Validation loss: 0.3448405755417688

F1 Score: 0.8841925591434728

D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:

UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn\_for)

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 5', max=158.0, style=ProgressStyle(
```

Epoch 5

Training loss: 0.18120793347494513

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value=''))))
```

Validation loss: 0.46844774058886934

F1 Score: 0.8689643477616807

D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:

UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.

'precision', 'predicted', average, warn\_for)

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 6', max=158.0, style=ProgressStyle(
```

Epoch 6

Training loss: 0.12532829719630978

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value=''))))
```

Validation loss: 0.513486019202641

F1 Score: 0.8811608118805235

```
D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:
UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels
with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 7', max=158.0, style=ProgressStyle(
```

Epoch 7

Training loss: 0.08559641675858558

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value=''))))
```

Validation loss: 0.5168239900044033

F1 Score: 0.8747685050375633

```
D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:
UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels
with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 8', max=158.0, style=ProgressStyle(
```

Epoch 8

Training loss: 0.07162901240436337

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value=''))))
```

Validation loss: 0.4332962855696678

F1 Score: 0.9068151300955131

```
D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:
UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels
with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 9', max=158.0, style=ProgressStyle(
```

Epoch 9

Training loss: 0.05525129703404028

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value='')))
```

```
Validation loss: 0.44452917150088717
```

```
F1 Score: 0.9045074396227775
```

```
D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:
```

```
UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels  
with no predicted samples.
```

```
    'precision', 'predicted', average, warn_for)
```

```
HBox(children=(FloatProgress(value=0.0, description='Epoch 10', max=158.0, style=ProgressStyle
```

```
Epoch 10
```

```
Training loss: 0.04888481118633777
```

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value='')))
```

```
Validation loss: 0.4590662739106587
```

```
F1 Score: 0.9074067359526623
```

```
D:\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135:
```

```
UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels  
with no predicted samples.
```

```
    'precision', 'predicted', average, warn_for)
```

## 10 Model Evaluation

```
[26]: avg_loss, preds, gts = evaluate(data_loader_val)  
      accuracy_per_class(preds, gts)
```

```
HBox(children=(FloatProgress(value=0.0, max=7.0), HTML(value='')))
```

```
Class: happy
```

```
Accuracy: 170/171
```

```
Class: not-relevant
```

```
Accuracy: 23/32
```

```
Class: angry
```

Accuracy: 6/9

Class: disgust

Accuracy: 0/1

Class: sad

Accuracy: 3/5

Class: surprise

Accuracy: 2/5

As it can be seen, the model did very good job in identifying happy sentences. However, in other cases its performance was significantly lower, such as 66% in angry and 40% in sad and surprise, respectively. Regarding disgust class, there was only one example for testing and the model could not predict it correctly.

That's it for this project! Thank you for following!